

The Production AI Readiness One-Pager

The 12 questions we ask before any engagement.

A diagnostic from TheFocus.AI for engineering leaders shipping LLM systems that have to work on Monday morning.

Why this one-pager exists

Most teams shipping "AI features" today are running a demo with extra steps. The model works in a notebook, works for the founder, works for the customer in the sales call. Then it ships, the contract closes, and the on-call engineer learns what the model does at 2 a.m. on a Tuesday when the API provider has a 30-second latency spike.

The 12 questions below are what we ask before we agree to help a team move from *"it works in the demo"* to *"it works in production, under load, with real users, and the bill is predictable."* If you can answer all 12 with a clear sentence, you do not need us. If three or more are fuzzy, you have a project — not a feature.

The 12 questions

PRODUCT SURFACE

1. **What user job does the AI do?** Name the single job-to-be-done in one sentence. If you need three, the surface is not ready.
2. **What is the failure mode the user sees?** When the model is wrong, slow, refuses, or unavailable — what does the screen look like and what does the user do next?
3. **What is the ground truth?** How do you know the answer was right? Who labels, who arbitrates, and how fast does that loop close?

SYSTEM ARCHITECTURE

4. **Where does the model live, and who pays for it?** API, fine-tune, self-hosted, on-device. Per-call cost, per-user cost, and who absorbs a 3× usage spike.
5. **What is the prompt boundary?** Where do user inputs end and your instructions begin? Have you red-teamed injection from the inputs you actually accept?
6. **What is in the context window, and who put it there?** RAG sources, retrieval policy, freshness, and what happens when retrieval returns nothing.

OPERATIONS

7. **What is the p95 latency budget?** Not the average. The 95th percentile a user will actually feel. Have you measured it under load with a real provider, not a local mock?
8. **What is the rate-limit and quota story?** Your provider's limits, your account's limits, and what happens when a single tenant burns the quota for the rest.
9. **How do you roll back a prompt?** Prompts are code. If a new prompt regresses a use case, what is the revert path, the audit trail, and the time-to-recover?

TRUST, SAFETY, AND THE BORING STUFF

10. **What data leaves your perimeter, and where does it land?** Provider retention, vendor sub-processors, contractual training opt-outs, and the legal review you actually completed (not the one in the slide deck).
11. **Who is accountable for a bad answer?** A named human, a Slack channel, an on-call rotation, a SLA. "The model" is not an accountable party.
12. **What does success look like in 90 days?** A metric, a baseline, a target, a review cadence, and a kill criterion if it does not work.

How to use this

Print it. Walk it through your team meeting. If three or more answers are fuzzy, you have homework before you ship — not a launch.

If you want a second pair of eyes on the homework, [we do that for a living](#).

TheFocus.AI — Building intelligent organizations. Production AI, not demoware.

[thefocus.ai](#) · [Subscribe to the weekly note](#)

v1 · 2026-05-25 · CTO draft — pending CEO review.